HireWise AI Incident Security Mapping — MITRE ATLAS Matrix

Consolidated 6-Entry Table Grouping Each Incident with Chosen Tactics, Techniques, and Procedures (TTPs)

Framework: MITRE ATLAS (v2026.09) | **System Evaluated:** HireWise AI LLM Recruitment & Candidate Ranking Pipeline | **Structure:** Exactly 6 Incident Entries (16 Mapped TTP Instances)
Evidence Legend: ■ **Green (#8ec843) — Directly Supported by Evidence:** Conclusively verified by logs, codebase artifacts, or observed outputs. ■ **Orange (#ffa500) — Implied / Suspected:** Architecturally inferred or hypothesized based on system behavior, not conclusively proven.

Incident Entry (HireWise Incident Report)	Chosen ATLAS TTP & Tactic	Evidence Status & Matrix Color	Supporting Case Evidence & Justification
Incident 1 **Manipulated CVs** HR received candidate applications containing hidden instructions in CV text. When ingested, HireWise's LLM ignored normal instructions and generated unrelated text in summaries.	**AML.T0051.001** **LLM Prompt Injection: Indirect** *Tactic:* Execution	■ **Green (#8ec843)** **Directly Supported by Evidence**	**Evidence:** CVs contained unusual hidden instructions; when processed, the assistant ignored its normal instructions and produced unrelated text in summaries. **Justification:** Malicious instructions arrived through an untrusted third-party data channel (the CV) ingested by the LLM, rather than from legitimate HR users. Downstream execution and alteration of output were directly observed.
	AML.T0065 **LLM Prompt Crafting** *Tactic:* AI Attack Adaptation	■ **Green (#8ec843)** **Directly Supported by Evidence**	**Evidence:** CV text contained deliberately authored prompt instructions designed to cause the assistant to drop its normal behavior. **Justification:** Adversary engineered and tuned prompt content tailored to the LLM summarizer so malicious directives would be executed. Even though the author's identity is unknown, the crafted payload was directly retrieved.
	AML.T0093 **Prompt Infiltration via Public-Facing Application** *Tactic:* Initial Access	▲ **Orange (#ffa500)** **Implied / Suspected**	**Evidence:** Manipulated CVs reached the HR team through the normal candidate application workflow. **Justification:** Suggests adversary placed prompts through an externally reachable intake channel for subsequent AI ingestion. Implied because the specific upload portal/mechanism is not explicitly detailed
	AML.T0068 **LLM Prompt Obfuscation** *Tactic:* Defense Evasion	▲ **Orange (#ffa500)** **Implied / Suspected**	**Evidence:** Adversarial instructions were described as 'hidden' within the CV body text. **Justification:** Hidden instructions suggest obfuscation techniques (e.g. tiny font, white text on white background, hidden metadata) to evade human recruiters and guardrails. Concealment method is not stated, so suspected.
	AML.T0054 **LLM Jailbreak** *Tactic:* Defense Evasion	▲ **Orange (#ffa500)** **Implied / Suspected**	**Evidence:** The assistant ignored its normal instructions and produced unrelated output text. **Justification:** Overriding intended behavior aligns with jailbreaking/guardrail circumvention, though the report does not formally distinguish between safety policy bypass and standard instruction hijacking.

Incident Entry (HireWise Incident Report)	Chosen ATLAS TTP & Tactic	Evidence Status & Matrix Color	Supporting Case Evidence & Justification
Incident 2 **Unauthorized Information Exposure** During testing, an HR employee asked the AI assistant to summarize a candidate. The summary leaked confidential internal notes about the candidate retrieved from the document database.	**AML.T0057** **LLM Data Leakage** *Tactic:* Exfiltration	■ **Green (#8ec843)** **Directly Supported by Evidence**	**Evidence:** When summarizing a candidate, the assistant included internal notes that were not meant to appear in the summary. **Justification:** The LLM disclosed confidential internal data directly in its output. Although observed during internal testing without hostile intent, the observed system behavior fits this exfiltration technique.
	AML.T0085.000 **Data from AI Services: RAG Databases** *Tactic:* Collection	▲ **Orange (#ffa500)** **Implied / Suspected**	**Evidence:** Internal notes stored in the internal document database appeared directly in the assistant's generated summary. **Justification:** Indicates the assistant dynamically retrieves records from a connected document store (RAG pattern) without sufficient access controls. The report does not name RAG explicitly, so this is inferred.
Incident 3 **Repeated API Queries** Security logs showed an employee account submitted hundreds of slightly modified prompts to the AI system asking for candidate rankings, scoring rationales, and response samples, saving outputs locally.	**AML.T0012** **Valid Accounts** *Tactic:* Initial Access	■ **Green (#8ec843)** **Directly Supported by Evidence**	**Evidence:** Security logs proved an employee account submitted hundreds of prompts to the AI system. **Justification:** Legitimate employee credentials were used to reach the system (insider abuse of a valid account).
	AML.T0040 **AI Model Inference API Access** *Tactic:* AI Model Access	■ **Green (#8ec843)** **Directly Supported by Evidence**	**Evidence:** Repeated API queries asking for candidate rankings, scoring explanations, and example responses. **Justification:** The account used the production inference API as its operational access path to probe the model and harvest its outputs.
	AML.T0127 **Data Staged** *Tactic:* Collection	■ **Green (#8ec843)** **Directly Supported by Evidence**	**Evidence:** The account saved the model outputs in local files on the employee's workstation. **Justification:** Model responses were gathered and consolidated in one staging location before being moved to the personal external project repository.
	AML.T0024.002 **Extract AI Model** *Tactic:* Exfiltration	■ **Green (#8ec843)** **Directly Supported by Evidence**	**Evidence:** Hundreds of slightly varied queries submitted to harvest candidate rankings and explanations. **Justification:** Repeatedly querying an inference API to collect outputs for mimic model training represents the query harvesting phase of model extraction.
	AML.T0063 **Discover AI Model Outputs** *Tactic:* Discovery	▲ **Orange (#ffa500)** **Implied / Suspected**	**Evidence:** Prompts requested scoring explanations and example responses, not just numerical rankings. **Justification:** Probing for scores and reasoning behind outputs suggests intentional reconnaissance of model reasoning logic to facilitate replication.

Incident Entry (HireWise Incident Report)	Chosen ATLAS TTP & Tactic	Evidence Status & Matrix Color	Supporting Case Evidence & Justification
Incident 4 **Suspicious External Model** Investigators found a personal project repository belonging to the employee containing collected AI outputs and a script using them to fine-tune a smaller LLM to reproduce HireWise's ranking behavior.	**AML.T0005.001** **Create Proxy AI Model: Train Proxy via Replication** *Tactic:* AI Attack Adaptation	■ **Green (#8ec843)** **Directly Supported by Evidence**	**Evidence:** Repository held a script that used collected HireWise outputs to fine-tune a smaller language model. **Justification:** Harvested model outputs were used as synthetic training labels/ground truth to fine-tune a surrogate proxy model mimicking target ranking behavior.
	AML.T0024.002 **Extract AI Model** *Tactic:* Exfiltration	■ **Green (#8ec843)** **Directly Supported by Evidence**	**Evidence:** Project repository explicitly documented its goal: to reproduce HireWise's candidate-ranking behavior. **Justification:** Realizes model extraction by distilling proprietary model functionality into an independent external asset.
Incident 5 **Possible Training-Data Manipulation** Several CVs were modified before upload with misleading qualifications and ranking instructions. Investigators confirmed no modifications occurred to underlying training datasets or model weights.	**AML.T0043.003** **Craft Adversarial Data: Manual Modification** *Tactic:* AI Attack Adaptation	■ **Green (#8ec843)** **Directly Supported by Evidence**	**Evidence:** Several CVs were deliberately modified before upload with misleading qualifications and ranking instructions. **Justification:** Manual modification of input data to push the model toward favorable output. Explicitly input-level manipulation; no training data or weights were altered.
	AML.T0051.001 **LLM Prompt Injection: Indirect** *Tactic:* Execution	■ **Green (#8ec843)** **Directly Supported by Evidence**	**Evidence:** CVs carried hidden instructions intended to influence the AI's candidate ranking calculations. **Justification:** Malicious instructions ingested through a third-party document channel (CV) to hijack the LLM scoring evaluation.
	AML.T0093 **Prompt Infiltration via Public-Facing Application** *Tactic:* Initial Access	▲ **Orange (#ffa500)** **Implied / Suspected**	**Evidence:** Manipulated CVs reached the HR team through standard application intake channels. **Justification:** Adversary introduced manipulated inputs via an external candidate intake channel. Intake infrastructure is inferred from recruitment workflows.
	AML.T0015 **Evade AI Model** *Tactic:* Defense Evasion	▲ **Orange (#ffa500)** **Implied / Suspected**	**Evidence:** CVs contained misleading qualifications and instructions intended to influence the AI's ranking. **Justification:** Intent was to make the ranking model misjudge candidate suitability. No evidence is given that any ranking was actually altered, so success is unconfirmed.
Incident 6 **Unconfirmed Impact** The replica model was shared with external users. While confirmed to be trained on HireWise outputs, investigators could not establish if it reproduced capabilities, leaked data, or enabled attacks.	**AML.T0048.004** **AI Intellectual Property Theft** *Tactic:* Impact	▲ **Orange (#ffa500)** **Implied / Suspected**	**Evidence:** The replica model was shared with other users; confirmed to be trained on HireWise outputs. **Justification:** A shared copy of proprietary ranking behavior implies possible IP theft. Investigators could not establish that it reproduced original capabilities, so impact is suspected only.